

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import tensorflow as tf
from keras.layers.normalization import BatchNormalization
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point
print(X_train[0])
```

[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	82	
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154		
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0															

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - Xmin)/(Xmax-Xmin) = X/255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

In [11]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10)
# ])
```

```

#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [12]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [13]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```

In [14]:

```

# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format

# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, step
s_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 1s 22us/step - loss: 0.6891 - accuracy: 0.8268 - val_loss: 0.3872 - val_accuracy: 0.8995
Epoch 2/20
60000/60000 [=====] - 1s 16us/step - loss: 0.3632 - accuracy: 0.9024 - val_loss: 0.3214 - val_accuracy: 0.9117
Epoch 3/20
60000/60000 [=====] - 1s 16us/step - loss: 0.3197 - accuracy: 0.9123 - val_loss: 0.2960 - val_accuracy: 0.9178
Epoch 4/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2996 - accuracy: 0.9169 - val_loss: 0.2844 - val_accuracy: 0.9205
Epoch 5/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2879 - accuracy: 0.9200 - val_loss: 0.2771 - val_accuracy: 0.9237
Epoch 6/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2796 - accuracy: 0.9218 - val_loss: 0.2752 - val_accuracy: 0.9224
Epoch 7/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2735 - accuracy: 0.9237 - val_loss: 0.2716 - val_accuracy: 0.9238
Epoch 8/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2691 - accuracy: 0.9243 - val_loss: 0.2678 - val_accuracy: 0.9252
Epoch 9/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2653 - accuracy: 0.9261 - val_loss: 0.2656 - val_accuracy: 0.9250
Epoch 10/20
60000/60000 [=====] - 1s 15us/step - loss: 0.2622 - accuracy: 0.9265 - val_loss: 0.2659 - val_accuracy: 0.9250
Epoch 11/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2595 - accuracy: 0.9279 - va

```

```

60000/60000 [=====] - 1s 16us/step - loss: 0.2554 - accuracy: 0.9289 - va
l_loss: 0.2635 - val_accuracy: 0.9265
Epoch 14/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2534 - accuracy: 0.9297 - va
l_loss: 0.2639 - val_accuracy: 0.9259
Epoch 15/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2520 - accuracy: 0.9303 - va
l_loss: 0.2613 - val_accuracy: 0.9275
Epoch 16/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2502 - accuracy: 0.9313 - va
l_loss: 0.2624 - val_accuracy: 0.9260
Epoch 17/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2493 - accuracy: 0.9302 - va
l_loss: 0.2637 - val_accuracy: 0.9263
Epoch 18/20
60000/60000 [=====] - 1s 16us/step - loss: 0.2482 - accuracy: 0.9312 - va
l_loss: 0.2625 - val_accuracy: 0.9280
Epoch 19/20
60000/60000 [=====] - 1s 17us/step - loss: 0.2470 - accuracy: 0.9321 - va
l_loss: 0.2632 - val_accuracy: 0.9263
Epoch 20/20
60000/60000 [=====] - 1s 18us/step - loss: 0.2459 - accuracy: 0.9322 - va
l_loss: 0.2625 - val_accuracy: 0.9268

```

In [15]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.26249462166130544
Test accuracy: 0.926800012588501





ARCHITECTURE 1= 784--612--324--10(2- Hidden layers)

A) MLP+ReLu+ADAM

In [16]:

```
model_relu = Sequential()
model_relu.add(Dense(612, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(324, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 612)	480420
dense_3 (Dense)	(None, 324)	198612
dense_4 (Dense)	(None, 10)	3250

Total params: 682,282
Trainable params: 682,282
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 28us/step - loss: 0.2157 - accuracy: 0.9359 - val_loss: 0.1121 - val_accuracy: 0.9635

Epoch 2/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0793 - accuracy: 0.9756 - val_loss: 0.0782 - val_accuracy: 0.9761

Epoch 3/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0477 - accuracy: 0.9849 - val_loss: 0.0680 - val_accuracy: 0.9777

Epoch 4/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0318 - accuracy: 0.9902 - val_loss: 0.0782 - val_accuracy: 0.9758

Epoch 5/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0246 - accuracy: 0.9918 - val_loss: 0.0696 - val_accuracy: 0.9778

Epoch 6/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0198 - accuracy: 0.9935 - val_loss: 0.0896 - val_accuracy: 0.9748

Epoch 7/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0152 - accuracy: 0.9950 - va


```

l_loss: 0.0785 - val_accuracy: 0.9775
Epoch 8/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0108 - accuracy: 0.9967 - va
l_loss: 0.0734 - val_accuracy: 0.9814
Epoch 9/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0180 - accuracy: 0.9940 - va
l_loss: 0.0939 - val_accuracy: 0.9755
Epoch 10/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0151 - accuracy: 0.9952 - va
l_loss: 0.1048 - val_accuracy: 0.9738
Epoch 11/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0144 - accuracy: 0.9952 - va
l_loss: 0.0825 - val_accuracy: 0.9802
Epoch 12/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0102 - accuracy: 0.9968 - va
l_loss: 0.0807 - val_accuracy: 0.9801
Epoch 13/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0095 - accuracy: 0.9965 - va
l_loss: 0.0902 - val_accuracy: 0.9805
Epoch 14/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0107 - accuracy: 0.9967 - va
l_loss: 0.0828 - val_accuracy: 0.9819
Epoch 15/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0094 - accuracy: 0.9969 - va
l_loss: 0.0836 - val_accuracy: 0.9826
Epoch 16/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0075 - accuracy: 0.9977 - va
l_loss: 0.1181 - val_accuracy: 0.9770
Epoch 17/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0116 - accuracy: 0.9962 - va
l_loss: 0.0840 - val_accuracy: 0.9826
Epoch 18/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0106 - accuracy: 0.9970 - va
l_loss: 0.1172 - val_accuracy: 0.9777
Epoch 19/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0071 - accuracy: 0.9976 - va
l_loss: 0.0965 - val_accuracy: 0.9820
Epoch 20/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0090 - accuracy: 0.9971 - va
l_loss: 0.0944 - val_accuracy: 0.9820

```

In [17]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.09443054281140775
Test accuracy: 0.9819999933242798

```



In [18]:

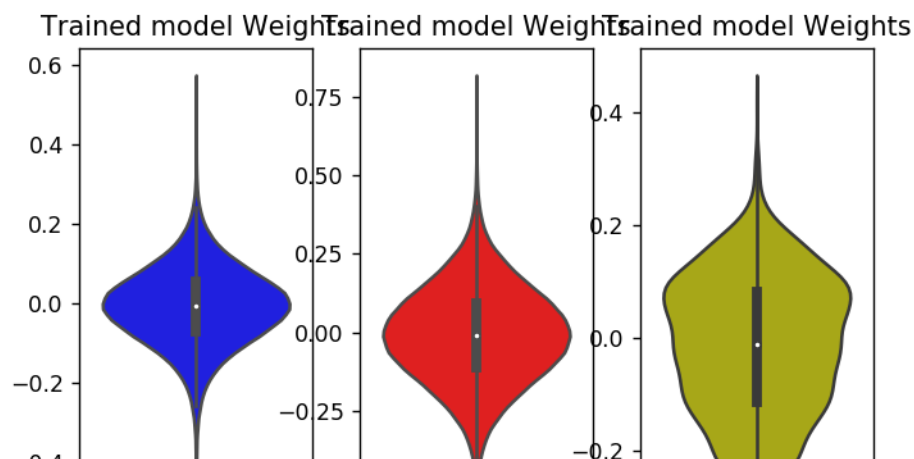
```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





ARCHITECTURE 1= 784--612--324--10(2- Hidden layers)

B) MLP+Dropout+ADAM

In [19]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(612, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(324, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 612)	480420
batch_normalization_1 (Batch Normalization)	(None, 612)	2448
dropout_1 (Dropout)	(None, 612)	0
dense_6 (Dense)	(None, 324)	198612
batch_normalization_2 (Batch Normalization)	(None, 324)	1296
dropout_2 (Dropout)	(None, 324)	0
dense_7 (Dense)	(None, 10)	3250
Total params: 686,026		
Trainable params: 684,154		
Non-trainable params: 1,872		

In [20]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20

```

60000/60000 [=====] - 3s 44us/step - loss: 0.4422 - accuracy: 0.8657 - va
l_loss: 0.1514 - val_accuracy: 0.9526
Epoch 2/20
60000/60000 [=====] - 2s 38us/step - loss: 0.2202 - accuracy: 0.9336 - va
l_loss: 0.1098 - val_accuracy: 0.9645
Epoch 3/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1737 - accuracy: 0.9467 - va
l_loss: 0.0985 - val_accuracy: 0.9685
Epoch 4/20
60000/60000 [=====] - 2s 38us/step - loss: 0.1525 - accuracy: 0.9529 - va
l_loss: 0.0879 - val_accuracy: 0.9716
Epoch 5/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1357 - accuracy: 0.9577 - va
l_loss: 0.0810 - val_accuracy: 0.9743
Epoch 6/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1251 - accuracy: 0.9608 - va
l_loss: 0.0771 - val_accuracy: 0.9766
Epoch 7/20
60000/60000 [=====] - 2s 38us/step - loss: 0.1153 - accuracy: 0.9640 - va
l_loss: 0.0731 - val_accuracy: 0.9770
Epoch 8/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1052 - accuracy: 0.9657 - va
l_loss: 0.0765 - val_accuracy: 0.9765
Epoch 9/20
60000/60000 [=====] - 2s 38us/step - loss: 0.1021 - accuracy: 0.9679 - va
l_loss: 0.0734 - val_accuracy: 0.9783
Epoch 10/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0951 - accuracy: 0.9706 - va
l_loss: 0.0692 - val_accuracy: 0.9790
Epoch 11/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0894 - accuracy: 0.9724 - va
l_loss: 0.0676 - val_accuracy: 0.9784
Epoch 12/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0828 - accuracy: 0.9741 - va
l_loss: 0.0598 - val_accuracy: 0.9821
Epoch 13/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0802 - accuracy: 0.9742 - va
l_loss: 0.0597 - val_accuracy: 0.9815
Epoch 14/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0769 - accuracy: 0.9757 - va
l_loss: 0.0636 - val_accuracy: 0.9802
Epoch 15/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0745 - accuracy: 0.9755 - va
l_loss: 0.0610 - val_accuracy: 0.9811
Epoch 16/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0682 - accuracy: 0.9787 - va
l_loss: 0.0608 - val_accuracy: 0.9802
Epoch 17/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0654 - accuracy: 0.9794 - va
l_loss: 0.0576 - val_accuracy: 0.9824
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0677 - accuracy: 0.9783 - va
l_loss: 0.0614 - val_accuracy: 0.9825
Epoch 19/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0639 - accuracy: 0.9798 - va
l_loss: 0.0565 - val_accuracy: 0.9813
Epoch 20/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0564 - accuracy: 0.9817 - va
l_loss: 0.0581 - val_accuracy: 0.9839

```

In [21]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

```

```


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05811911709617998
Test accuracy: 0.9839000105857849



In [22]:

```

w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



ARCHITECTURE 2= 784--438--276--157--10(3- Hidden layers)

A) MLP+ReLu+ADAM

In [23]:

```
model_relu = Sequential()
model_relu.add(Dense(438, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(276, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(157, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 438)	343830
dense_9 (Dense)	(None, 276)	121164
dense_10 (Dense)	(None, 157)	43489
dense_11 (Dense)	(None, 10)	1580

Total params: 510,063

Trainable params: 510,063

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 30us/step - loss: 0.2358 - accuracy: 0.9292 - val_loss: 0.1033 - val_accuracy: 0.9669

Epoch 2/20

```

60000/60000 [=====] - 2s 28us/step - loss: 0.0864 - accuracy: 0.9733 - va
l_loss: 0.0997 - val_accuracy: 0.9700
Epoch 3/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0558 - accuracy: 0.9829 - va
l_loss: 0.0973 - val_accuracy: 0.9686
Epoch 4/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0394 - accuracy: 0.9875 - va
l_loss: 0.0900 - val_accuracy: 0.9736
Epoch 5/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0326 - accuracy: 0.9893 - va
l_loss: 0.0941 - val_accuracy: 0.9743
Epoch 6/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0300 - accuracy: 0.9898 - va
l_loss: 0.0912 - val_accuracy: 0.9716
Epoch 7/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0204 - accuracy: 0.9934 - va
l_loss: 0.0793 - val_accuracy: 0.9794
Epoch 8/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0203 - accuracy: 0.9933 - va
l_loss: 0.0831 - val_accuracy: 0.9792
Epoch 9/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0193 - accuracy: 0.9935 - va
l_loss: 0.0822 - val_accuracy: 0.9794
Epoch 10/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0158 - accuracy: 0.9943 - va
l_loss: 0.1041 - val_accuracy: 0.9764
Epoch 11/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0155 - accuracy: 0.9952 - va
l_loss: 0.1100 - val_accuracy: 0.9760
Epoch 12/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0182 - accuracy: 0.9941 - va
l_loss: 0.1010 - val_accuracy: 0.9794
Epoch 13/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0145 - accuracy: 0.9949 - va
l_loss: 0.1080 - val_accuracy: 0.9768
Epoch 14/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0140 - accuracy: 0.9951 - va
l_loss: 0.0903 - val_accuracy: 0.9819
Epoch 15/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0107 - accuracy: 0.9967 - va
l_loss: 0.1016 - val_accuracy: 0.9823
Epoch 16/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0112 - accuracy: 0.9964 - va
l_loss: 0.0986 - val_accuracy: 0.9800
Epoch 17/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0113 - accuracy: 0.9963 - va
l_loss: 0.0910 - val_accuracy: 0.9808
Epoch 18/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0101 - accuracy: 0.9965 - va
l_loss: 0.1061 - val_accuracy: 0.9797
Epoch 19/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0124 - accuracy: 0.9964 - va
l_loss: 0.1029 - val_accuracy: 0.9783
Epoch 20/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0096 - accuracy: 0.9970 - va
l_loss: 0.0927 - val_accuracy: 0.9814

```

In [24]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val loss : validation loss


```

```
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09267379564880103
 Test accuracy: 0.9814000129699707



In [25]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

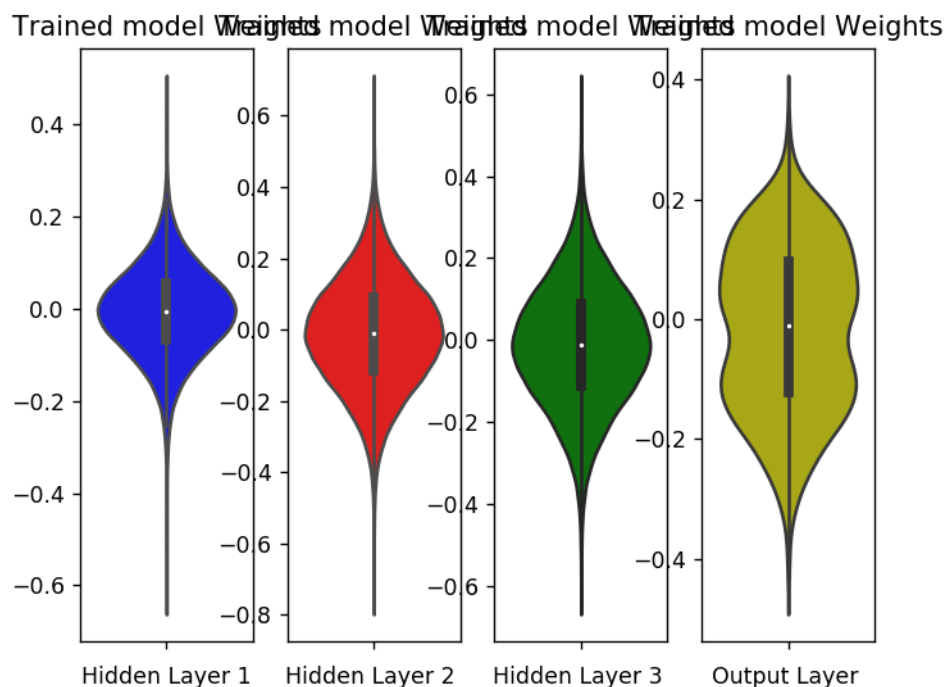
plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
```



```
plt.show()
```



ARCHITECTURE 2= 784--438--276--157--10(3- Hidden layers)

B) MLP+Dropout+ADAM

In [26]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(438, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(276, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(157, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 438)	343830
batch normalization 3 (Batch Normalization)	(None, 438)	1752

dropout_3 (Dropout)	(None, 438)	0
dense_13 (Dense)	(None, 276)	121164
batch_normalization_4 (Batch Normalization)	(None, 276)	1104
dropout_4 (Dropout)	(None, 276)	0
dense_14 (Dense)	(None, 157)	43489
batch_normalization_5 (Batch Normalization)	(None, 157)	628
dropout_5 (Dropout)	(None, 157)	0
dense_15 (Dense)	(None, 10)	1580
=====		
Total params: 513,547		
Trainable params: 511,805		
Non-trainable params: 1,742		

In [27]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 52us/step - loss: 0.5875 - accuracy: 0.8211 - val_loss: 0.1787 - val_accuracy: 0.9440
Epoch 2/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2806 - accuracy: 0.9160 - val_loss: 0.1358 - val_accuracy: 0.9572
Epoch 3/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2264 - accuracy: 0.9320 - val_loss: 0.1200 - val_accuracy: 0.9633
Epoch 4/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1899 - accuracy: 0.9431 - val_loss: 0.1012 - val_accuracy: 0.9704
Epoch 5/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1700 - accuracy: 0.9491 - val_loss: 0.1023 - val_accuracy: 0.9704
Epoch 6/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1565 - accuracy: 0.9526 - val_loss: 0.0946 - val_accuracy: 0.9723
Epoch 7/20
60000/60000 [=====] - 3s 45us/step - loss: 0.1473 - accuracy: 0.9553 - val_loss: 0.0823 - val_accuracy: 0.9742
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1329 - accuracy: 0.9597 - val_loss: 0.0851 - val_accuracy: 0.9756
Epoch 9/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1276 - accuracy: 0.9611 - val_loss: 0.0741 - val_accuracy: 0.9780
Epoch 10/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1190 - accuracy: 0.9638 - val_loss: 0.0814 - val_accuracy: 0.9759
Epoch 11/20
60000/60000 [=====] - 3s 45us/step - loss: 0.1141 - accuracy: 0.9650 - val_loss: 0.0710 - val_accuracy: 0.9794
Epoch 12/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1088 - accuracy: 0.9675 - val_loss: 0.0744 - val_accuracy: 0.9785
Epoch 13/20
60000/60000 [=====] - 3s 44us/step - loss: 0.1041 - accuracy: 0.9686 - val_loss: 0.0707 - val_accuracy: 0.9790
Epoch 14/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0982 - accuracy: 0.9712 - val_loss: 0.0667 - val_accuracy: 0.9800
Epoch 15/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0910 - accuracy: 0.9724 - val_loss: 0.0730 - val_accuracy: 0.9782
```

```

Epoch 16/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0946 - accuracy: 0.9720 - va
l_loss: 0.0683 - val_accuracy: 0.9796
Epoch 17/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0854 - accuracy: 0.9738 - va
l_loss: 0.0734 - val_accuracy: 0.9792
Epoch 18/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0864 - accuracy: 0.9740 - va
l_loss: 0.0646 - val_accuracy: 0.9817
Epoch 19/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0850 - accuracy: 0.9746 - va
l_loss: 0.0668 - val_accuracy: 0.9816
Epoch 20/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0772 - accuracy: 0.9761 - va
l_loss: 0.0666 - val_accuracy: 0.9795

```

In [28]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0666387884610449
Test accuracy: 0.9794999957084656



2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0
epoch

In [29]:

```
w_after = model_drop.get_weights()

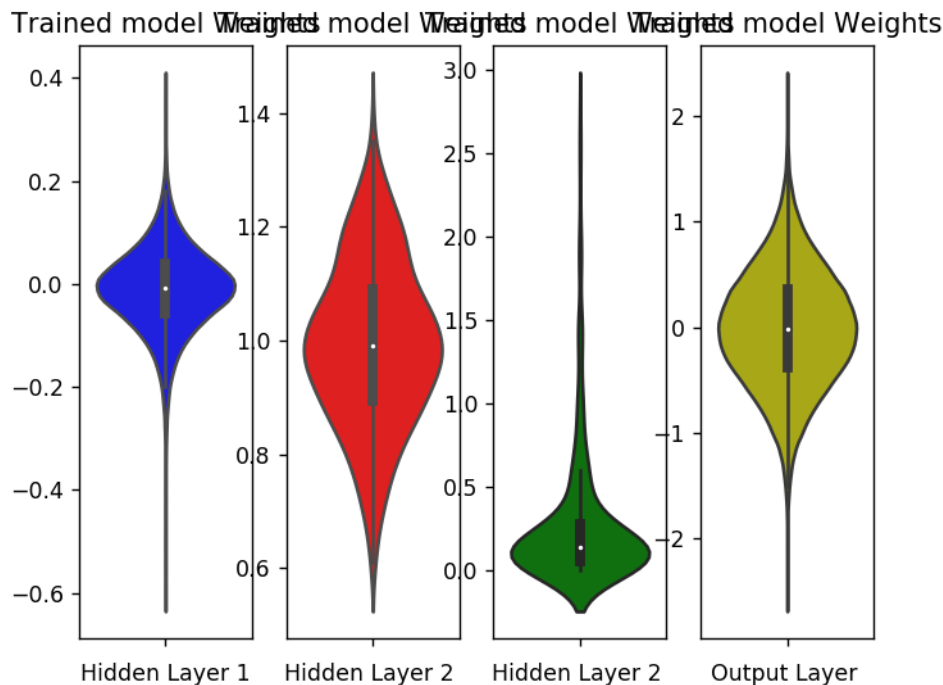
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



ARCHITECTURE 3= 784--551--447--331--236--121--10(5- Hidden layers)

A) MLP+ReLu+ADAM

In [30]:

```
model_relu = Sequential()
model_relu.add(Dense(551, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(447, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(331, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(236, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(121, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 551)	432535
dense_17 (Dense)	(None, 447)	246744
dense_18 (Dense)	(None, 331)	148288
dense_19 (Dense)	(None, 236)	78352
dense_20 (Dense)	(None, 121)	28677
dense_21 (Dense)	(None, 10)	1220

Total params: 935,816
Trainable params: 935,816
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 37us/step - loss: 0.2758 - accuracy: 0.9219 - val_loss: 0.1346 - val_accuracy: 0.9587

Epoch 2/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0971 - accuracy: 0.9699 - val_loss: 0.1024 - val_accuracy: 0.9675

Epoch 3/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0677 - accuracy: 0.9788 - val_loss: 0.1066 - val_accuracy: 0.9661

Epoch 4/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0525 - accuracy: 0.9829 - val_loss: 0.1103 - val_accuracy: 0.9695

Epoch 5/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0422 - accuracy: 0.9867 - val_loss: 0.0919 - val_accuracy: 0.9744

Epoch 6/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0366 - accuracy: 0.9878 - val_loss: 0.1137 - val_accuracy: 0.9713

Epoch 7/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0381 - accuracy: 0.9879 - val_loss: 0.1059 - val_accuracy: 0.9716

Epoch 8/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0298 - accuracy: 0.9902 - val_loss: 0.1128 - val_accuracy: 0.9710

Epoch 9/20

60000/60000 [=====] - 2s 33us/step - loss: 0.0297 - accuracy: 0.9906 - val_loss: 0.1026 - val_accuracy: 0.9745

Epoch 10/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0269 - accuracy: 0.9912 - val_loss: 0.1038 - val_accuracy: 0.9738

Epoch 11/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0241 - accuracy: 0.9927 - val_loss: 0.1000 - val_accuracy: 0.9750

```

l_loss: 0.0901 - val_accuracy: 0.9800
Epoch 12/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0221 - accuracy: 0.9930 - va
l_loss: 0.1011 - val_accuracy: 0.9771
Epoch 13/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0200 - accuracy: 0.9937 - va
l_loss: 0.1163 - val_accuracy: 0.9735
Epoch 14/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0169 - accuracy: 0.9946 - va
l_loss: 0.0861 - val_accuracy: 0.9819
Epoch 15/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0193 - accuracy: 0.9947 - va
l_loss: 0.0982 - val_accuracy: 0.9776
Epoch 16/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0201 - accuracy: 0.9941 - va
l_loss: 0.0883 - val_accuracy: 0.9793
Epoch 17/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0126 - accuracy: 0.9964 - va
l_loss: 0.1204 - val_accuracy: 0.9753
Epoch 18/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0184 - accuracy: 0.9946 - va
l_loss: 0.1000 - val_accuracy: 0.9772
Epoch 19/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0136 - accuracy: 0.9960 - va
l_loss: 0.0934 - val_accuracy: 0.9814
Epoch 20/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0155 - accuracy: 0.9956 - va
l_loss: 0.0935 - val_accuracy: 0.9804

```

In [31]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09345374069254558
Test accuracy: 0.980400025844574





In [32]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained\n")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

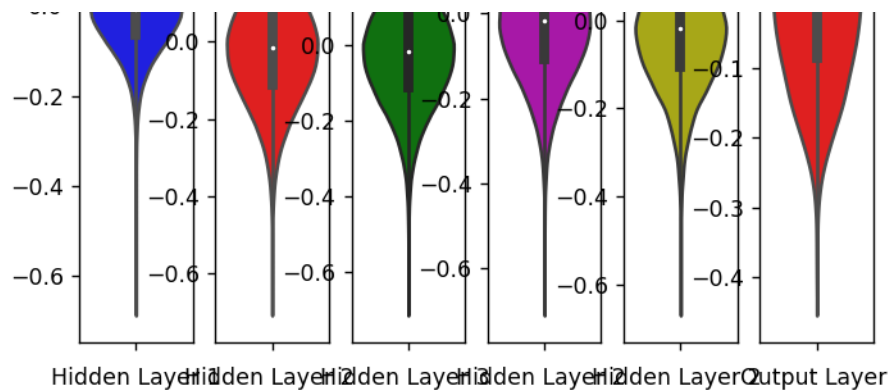
plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='m')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='y')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='r')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights, Hidden Layer 1 Weights, Hidden Layer 2 Weights, Hidden Layer 3 Weights, Hidden Layer 2 Weights, Hidden Layer 2 Weights, Output Layer Weights





ARCHITECTURE 3= 784--551--447--331--236--121--10(5- Hidden layers)

B) MLP+Dropout+ADAM

In [33]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(551, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(447, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(331, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(236, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(121, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 551)	432535
batch_normalization_6 (Batch Normalization)	(None, 551)	2204
dropout_6 (Dropout)	(None, 551)	0
dense_23 (Dense)	(None, 447)	246744

batch_normalization_7 (Batch Normalization)	(None, 447)	1788
dropout_7 (Dropout)	(None, 447)	0
dense_24 (Dense)	(None, 331)	148288
batch_normalization_8 (Batch Normalization)	(None, 331)	1324
dropout_8 (Dropout)	(None, 331)	0
dense_25 (Dense)	(None, 236)	78352
batch_normalization_9 (Batch Normalization)	(None, 236)	944
dropout_9 (Dropout)	(None, 236)	0
dense_26 (Dense)	(None, 121)	28677
batch_normalization_10 (Batch Normalization)	(None, 121)	484
dropout_10 (Dropout)	(None, 121)	0
dense_27 (Dense)	(None, 10)	1220
=====		
Total params: 942,560		
Trainable params: 939,188		
Non-trainable params: 3,372		

In [34]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 76us/step - loss: 0.8836 - accuracy: 0.7288 - val_loss: 0.2308 - val_accuracy: 0.9279
Epoch 2/20
60000/60000 [=====] - 4s 61us/step - loss: 0.3485 - accuracy: 0.8990 - val_loss: 0.1723 - val_accuracy: 0.9492
Epoch 3/20
60000/60000 [=====] - 4s 62us/step - loss: 0.2714 - accuracy: 0.9225 - val_loss: 0.1403 - val_accuracy: 0.9592
Epoch 4/20
60000/60000 [=====] - 4s 61us/step - loss: 0.2290 - accuracy: 0.9353 - val_loss: 0.1156 - val_accuracy: 0.9666
Epoch 5/20
60000/60000 [=====] - 4s 62us/step - loss: 0.2026 - accuracy: 0.9427 - val_loss: 0.1083 - val_accuracy: 0.9691
Epoch 6/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1872 - accuracy: 0.9463 - val_loss: 0.1063 - val_accuracy: 0.9708
Epoch 7/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1725 - accuracy: 0.9505 - val_loss: 0.1015 - val_accuracy: 0.9709
Epoch 8/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1624 - accuracy: 0.9543 - val_loss: 0.0982 - val_accuracy: 0.9715
Epoch 9/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1524 - accuracy: 0.9571 - val_loss: 0.0867 - val_accuracy: 0.9752
Epoch 10/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1413 - accuracy: 0.9604 - val_loss: 0.0796 - val_accuracy: 0.9793
Epoch 11/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1354 - accuracy: 0.9614 - val_loss: 0.0797 - val_accuracy: 0.9774
Epoch 12/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1296 - accuracy: 0.9629 - val_loss: 0.0803 - val_accuracy: 0.9782
Epoch 13/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1253 - accuracy: 0.9645 - va
```

```

l_loss: 0.0749 - val_accuracy: 0.9798
Epoch 14/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1193 - accuracy: 0.9658 - va
l_loss: 0.0813 - val_accuracy: 0.9787
Epoch 15/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1150 - accuracy: 0.9681 - va
l_loss: 0.0812 - val_accuracy: 0.9786
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1129 - accuracy: 0.9677 - va
l_loss: 0.0805 - val_accuracy: 0.9786
Epoch 17/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1098 - accuracy: 0.9683 - va
l_loss: 0.0773 - val_accuracy: 0.9792
Epoch 18/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1032 - accuracy: 0.9712 - va
l_loss: 0.0674 - val_accuracy: 0.9822
Epoch 19/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0965 - accuracy: 0.9728 - va
l_loss: 0.0708 - val_accuracy: 0.9811
Epoch 20/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0971 - accuracy: 0.9725 - va
l_loss: 0.0690 - val_accuracy: 0.9806

```

In [35]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06895258962376974
Test accuracy: 0.9805999994277954





In [36]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[6].flatten().reshape(-1,1)
h4_w = w_after[10].flatten().reshape(-1,1)
h5_w = w_after[14].flatten().reshape(-1,1)
out_w = w_after[18].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='m')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights, Hidden Layer 1 Weights, Hidden Layer 2 Weights, Hidden Layer 3 Weights, Hidden Layer 4 Weights, Hidden Layer 5 Weights, Output Layer Weights





In []: